

CS5223
Distributed Systems

Lecture 2: Communication

Instructor: YU Haifeng

Today's Roadmap

- Chapter 4 of textbook
- Review (Overview) of the OSI model
- Remote Procedure Call / Remote Method Invocation
 - If you don't know Java, read some tutorials on Java (e.g., "Java in a Nutshell"
https://www.amazon.com/Java-Nutshell-Desktop-Quick-Reference-dp-1098131002/dp/1098131002/ref=dp_ob_title_bk)
 - Should be quite easy if you already know C++
- Multicast
- Gossiping

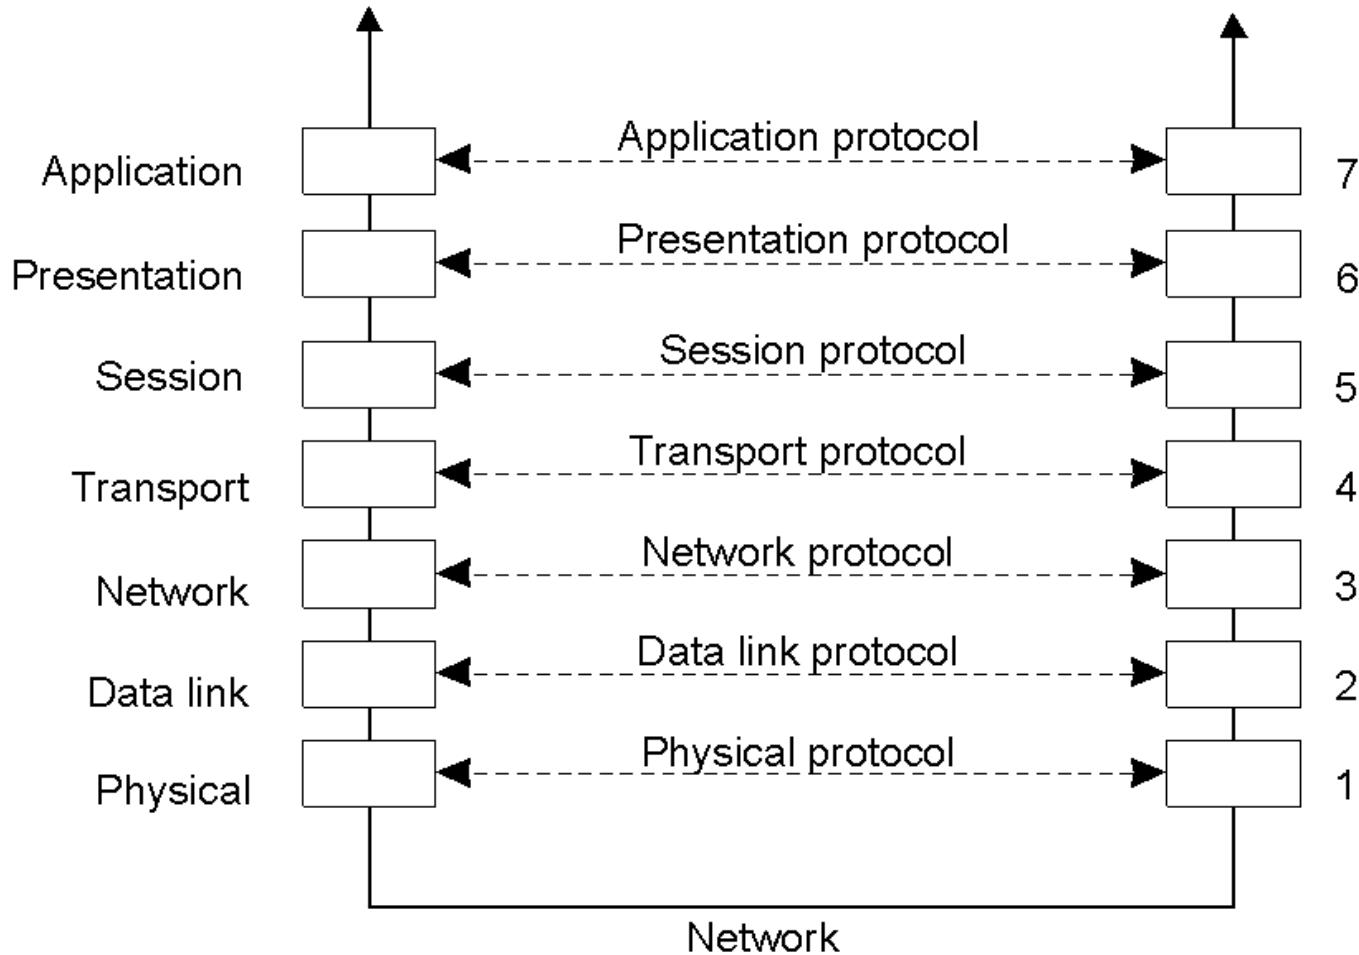
Motivation

- Distributed systems by definition need communication
- Networking:
 - Aims at providing delivery of data
 - Does not care about the data content
- Distributed systems:
 - The layer above networking
 - Does not care how the data is delivered – **sometimes** treat the network as a black-box
- Pros and cons of treating network as a black box?
 - Example?

distributed
systems

networking

Overview (Review) of OSI Model



Open Systems Interconnection Reference Model

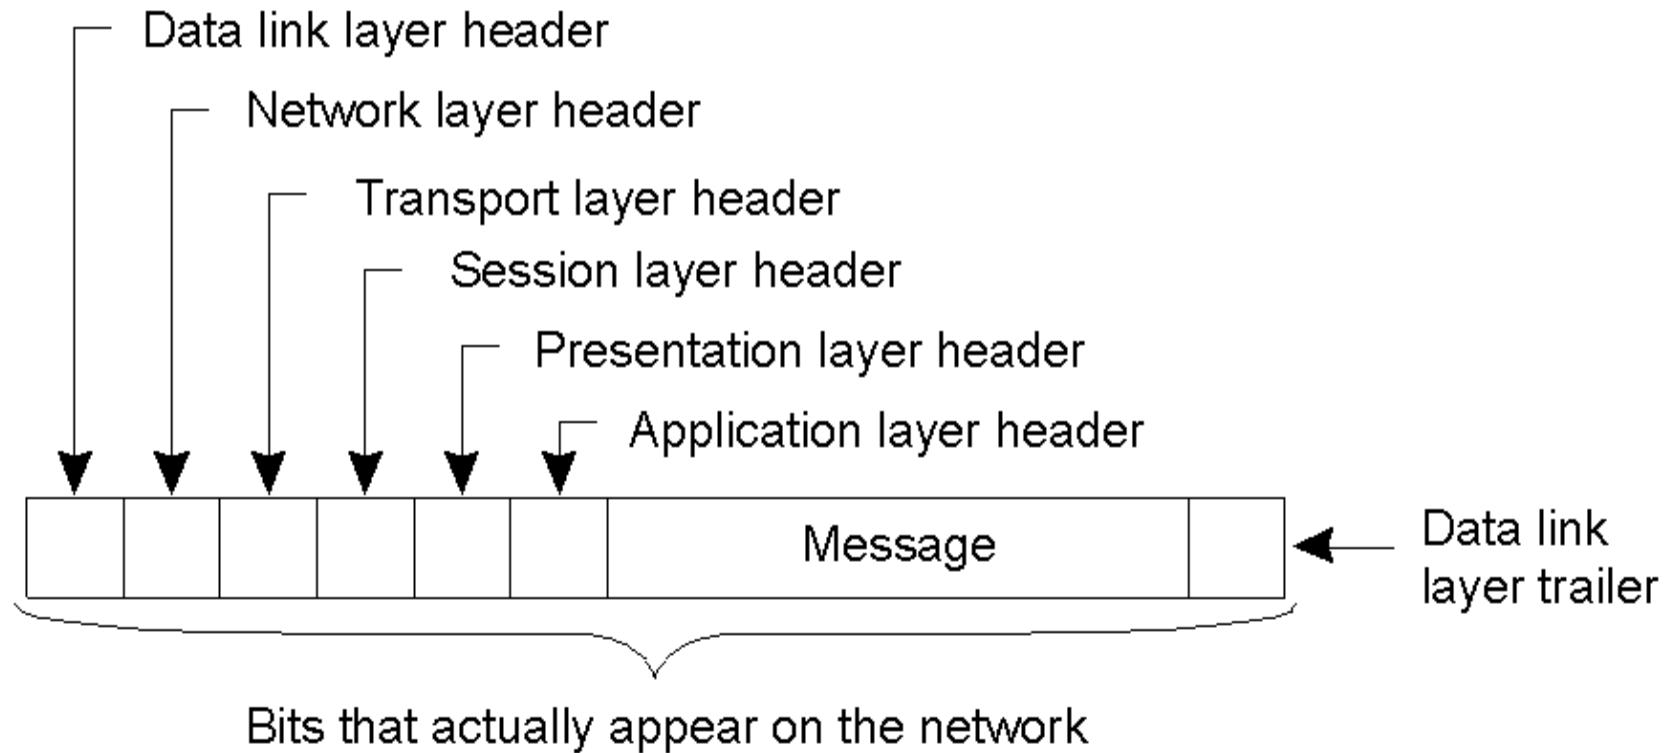
More Details on Each Layer

- Physical layer
 - Transmission of bits between sender and receiver
- Data link layer
 - Transmission of frames (series of bits) with error correction and detection
- Network layer
 - Describes how packets are *routed along multiple hops*
 - Prominent example: IP (Internet protocol)

Transport Layer

- Transport layer provides the actual communication facilities for most distributed systems
- Standard Internet transport-layer protocols
 - TCP (transmission control protocol): Reliable, in-order, with flow control, stateful
 - UDP (universal datagram protocol): Unreliable (best-effort) datagram connectionless communication

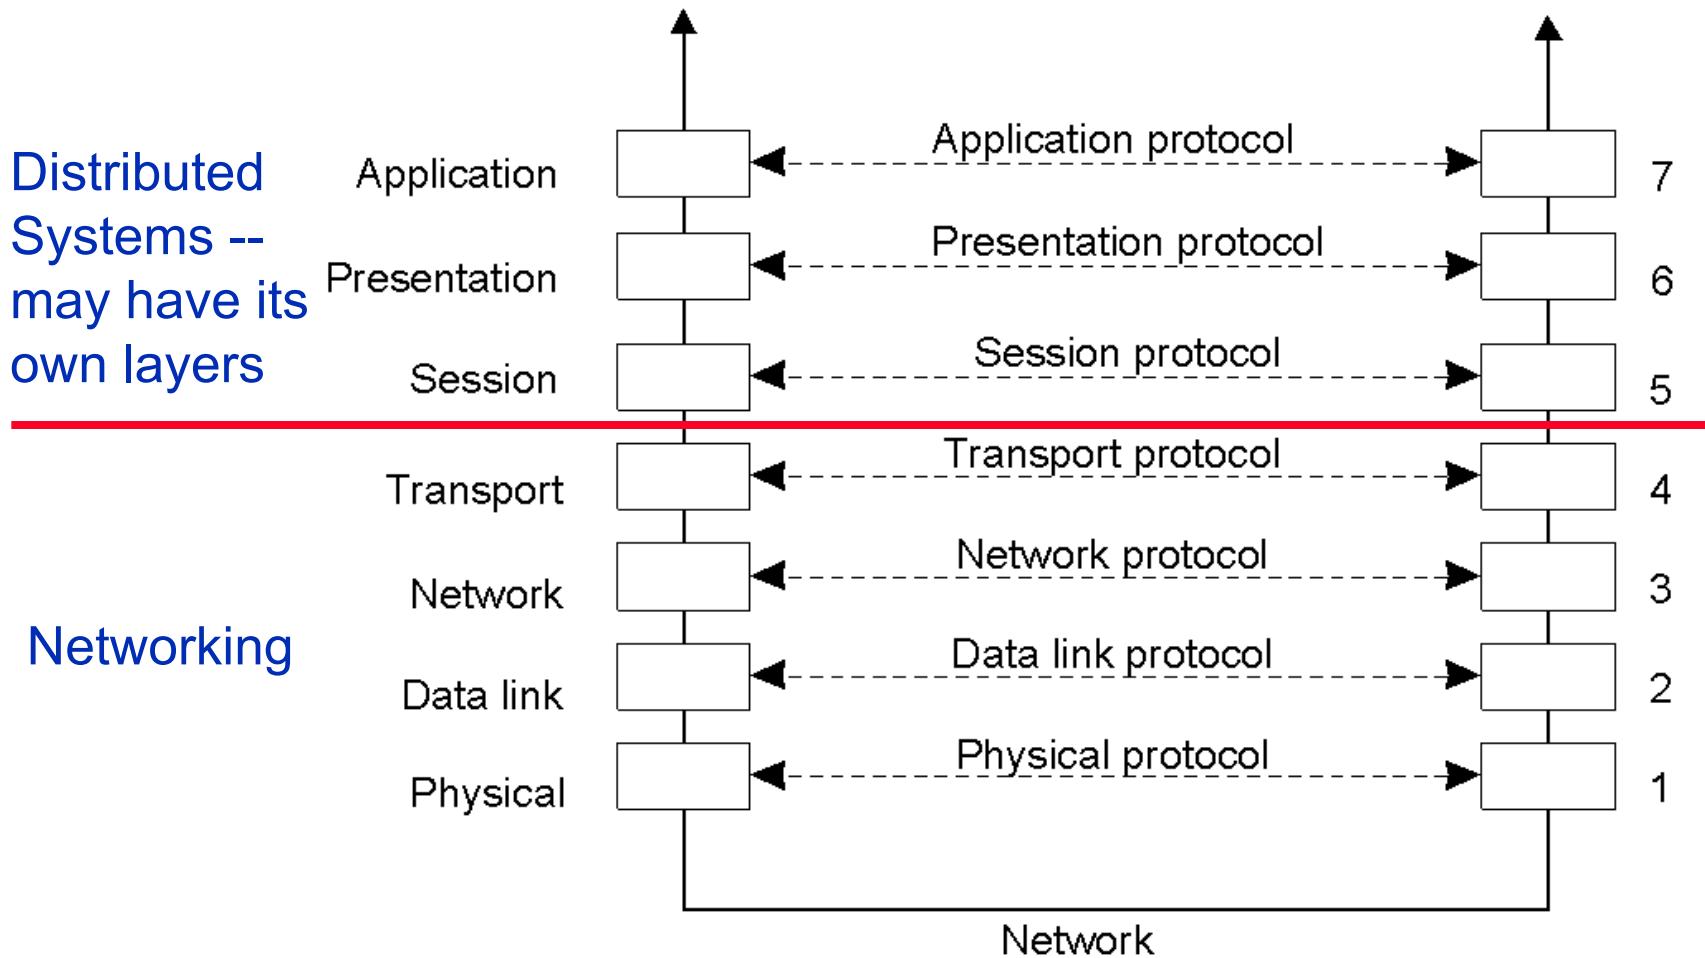
Messages in the OSI Model



OSI Model - Continued

- The OSI model does not dictate how you implement each layer
- Nothing fancy – just the application of layered modular design principle
 - Neither the only model nor the best model (no single best model)
 - But it is widely accepted – you want to speak the same language as other people
- Should NOT view it as cast into stone: **The layered modular design principle is more important than the model itself**
- OSI aims for generality, and one can also customize it in different context
 - Why should we customize it?
 - Example?

The concepts of Application, Presentation, and Session layers are not widely used.



Motivation for More Powerful Communication Mechanisms

- Many distributed systems directly use TCP/UDP
 - Prominent examples: Distributed games, BitTorrent
 - In theory, all distributed systems can be built by directly using TCP/UDP
- But many distributed systems share common communication pattern / requirements
 - Middleware / libraries / tools / modules to simplify their design
 - Again, just a software engineering principle (what principle?)

More Powerful Communication Mechanisms

- Remote Procedure Call / Remote Method Invocation
- Multicast
- Gossiping

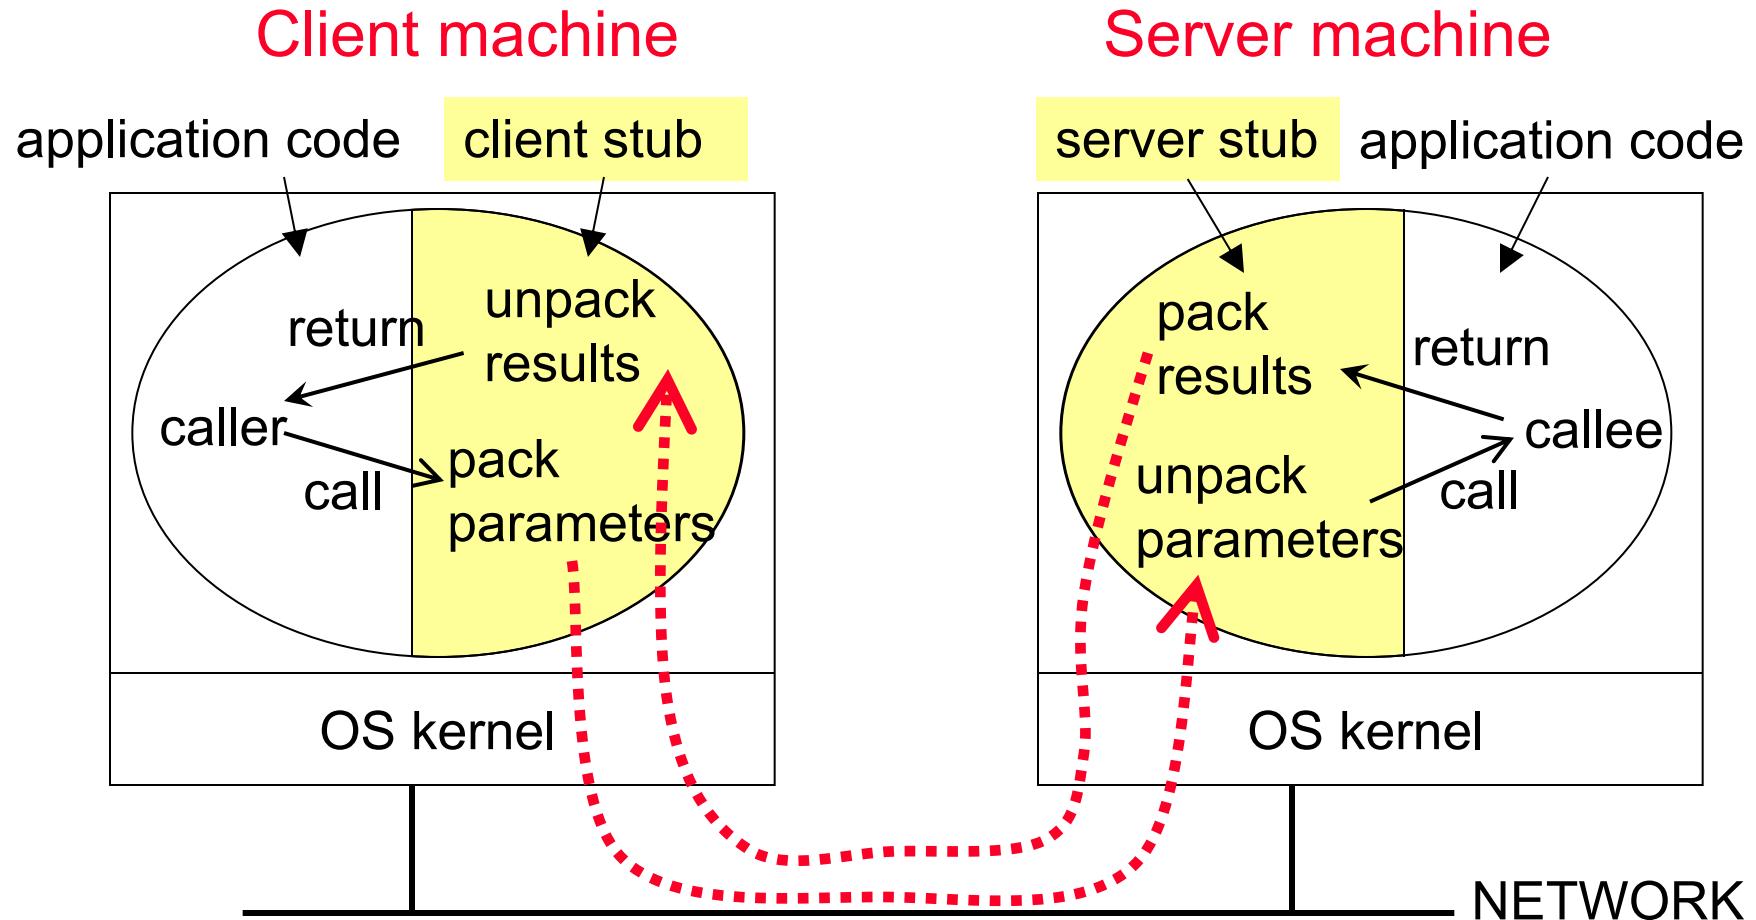
Motivation: Remote Procedure Call

- Many distributed systems involve the interactions between a client and a server
 - Example: Clients wants to read the content of a file on a remote machine
 - So common that we want to develop a common middleware / library / toolkit to simplify their designs
- One approach (not the only approach): Remote procedure call
 - Example: **read(fileid, buffer, numbytes)**

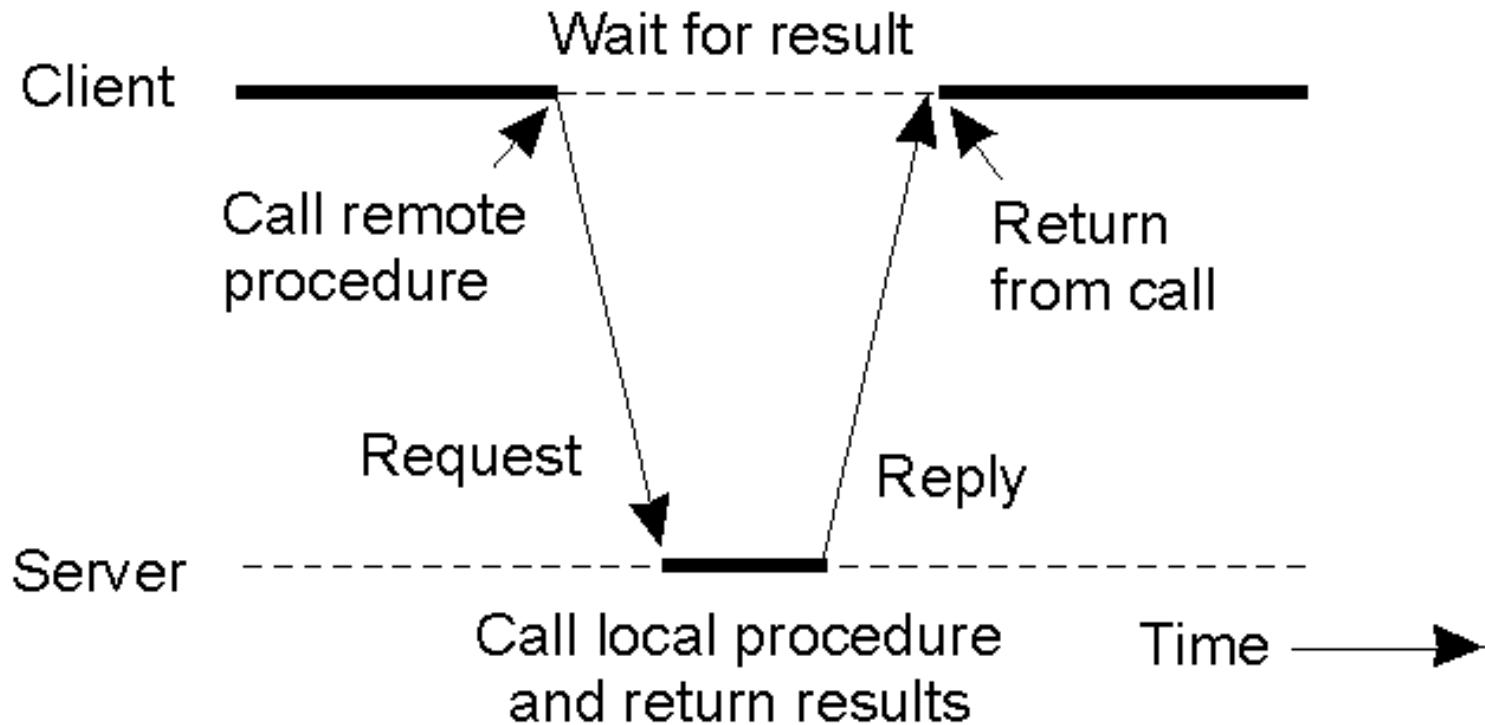
Motivation: Remote Procedure Call

- Sending the parameters to the remote machine and then wait for a reply
 - Simple (and potentially efficient) to implement
 - Can do so using a general middleware that work for all remote procedure calls
- Why RPCs are good
 - Application developers are familiar with procedure calls
 - RPCs operate in isolation – modular design
- History has proved the huge success of RPC
 - Java and C# have RPC support built into the language

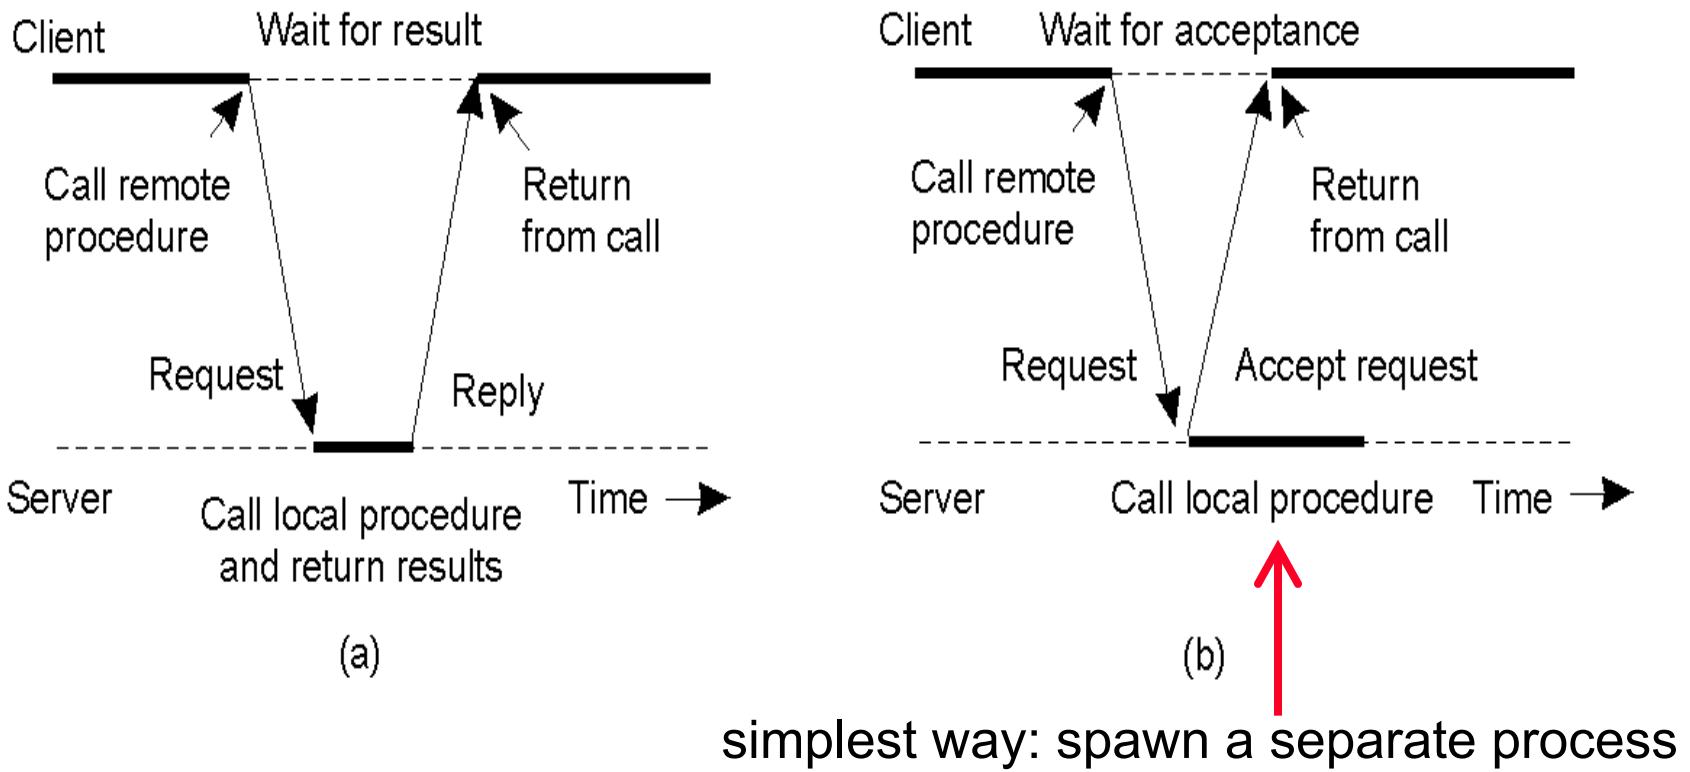
Typical Architecture for Enabling RPC



RPCs are by definition blocking



Making RPCs Non-blocking



- Non-blocking RPCs depart from procedure call semantics
- If you use a lot of non-blocking RPCs, re-think whether you indeed want to use RPC

RMI: A Case Study for RPC

- RMI is just object-oriented RPC
 - Each function (i.e., method) in Java belongs to some object
 - `read(file, buffer, numbytes)` ⇒ `file.read(buffer, numbytes)`
- **file** is a reference to a “remote object”
- We will focus on RMI as an example from now on

Java Remote Objects

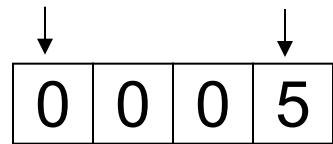
- What is needed in a remote object reference
 - Need to ensure uniqueness over space – no other object has the same “name”
 - Need to ensure uniqueness over time – no other objects should have the same “name” in the past or in the future

IP addr	port num	creation time	obj num	methods
---------	----------	---------------	---------	---------

Marshaling: Primitive Data Types

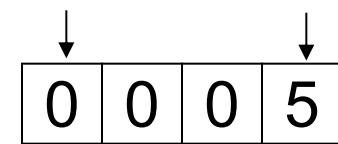
- Marshaling: Client needs to send the parameters of the function call to the server
 - Wrap all parameters into a message and send the message
- Network transmits raw bytes: Client and server may have different representation of data

most least
significant significant



4-byte integer on Pentium

least most
significant significant



4-byte integer on SPARC

Marshaling: Java (Non-remote) Objects

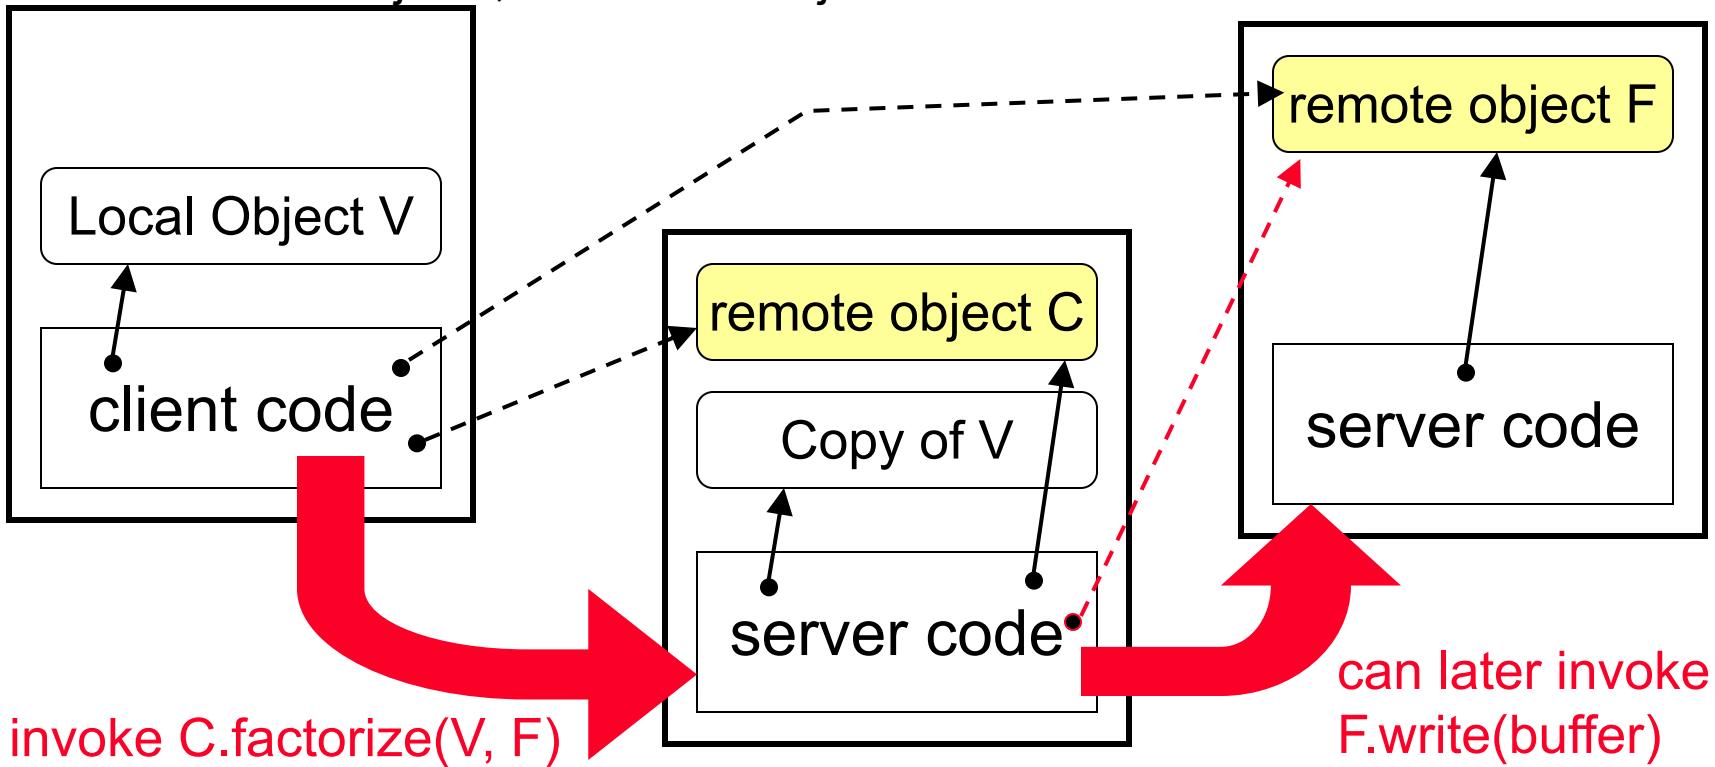
- How to marshal complex data types (i.e., objects)
 - Example: A Vector of integers
 - Java object **serialization**: convert a Java object to a bit string (for sending over the network or storing on disk)
- Java serialization extremely useful:
 - How to store a binary tree?
- What about a vector of vectors?
 - When serializing an obj, all obj it references are serialized as well (and this goes on...how?)
 - **Understand the object you are serializing**

Marshaling: Java (Non-remote) Objects

- Unmarshaling on server: Use **de-serialization** and convert the bit stream back to object
 - A new object will be created on server
- Use of object serialization/de-serialization ⇒
Semantics of RMI is different from local method invocation
 - Server will work on the “cloned” object
 - The “clone” can then start differing from the original....

Marshaling: Java Remote Objects

- Remote objects are passed by reference when used as parameter
 - `C.factorize(V, F)`
 - C is a ComputeEngine object; V is a Vector of integers, F is output file
 - C and F are remote objects, V is a local object



Marshaling: Java Remote Objects

- Remote objects are always “interfaces”
- The remote object may originates from the client or some other server
 - Enables complex interactions (such as callbacks) – how?

Result Marshaling/Unmarshaling

- Exactly the same as parameter marshaling / unmarshaling

Naming in Java RMI

- Obtain a reference of a remote object for the first time
 - Server machine runs rmiregistry (a separate process from the Java process)
`rmiregistry [port]`
 - Java prog on the server invokes
`Naming.rebind("objname", remote_obj);`
 - rmiregistry will remember this name binding
 - Java prog on the client invokes
`obj = Naming.lookup("//ip_addr:port/objname");`
 - rmiregistry will return the remote obj
- Lookup itself is like an RMI

Downloading of Classes

- Two scenario:
 - Receive a marshaled object, but don't have the code (i.e., class file) for that object
 - Receive a reference of a remote object, but don't have the code (i.e., class file) for the object
- Java will automatically “download” the class file from the machine where you get the object
 - Advantage: Allow automatic object upgrade
 - Disadvantage: **If you already have a class file for that object, the local class file will be used – possibility of inconsistency**

When we should NOT use RMI/RPC

- RPC aims for client/server communication:
 - Not suitable for multicast/broadcast
- Performance: Why RPC is not used for web browsing?
 - Sub-question: What exactly are we using for web browsing?
 - All about tradeoffs
- Blocking: A client cannot invoke RPCs in parallel
 - Non-blocking RPC solves the problem, but at the cost of design complexity – sometimes beats the exact purpose of RPC!

When we should NOT use RMI/RPC

- Failures: Perhaps the biggest vulnerability
 - Concept of RPC comes from local procedure call and intends to preserve its behavior
 - What if the server fails during execution?
 - Client has no idea whether the procedure has been invoked or not ...example?
 - To ensure exactly-once semantics, careful /complex logging needs to be done
- Should we hide such failure or expose it?
 - Again, it is about tradeoffs

How to make up your mind?

Rule of thumb: Use RPC if the communication pattern is simple, if you do not expect/worry much about failures, and if performance overheads are acceptable.

History Readings (Non-compulsory)

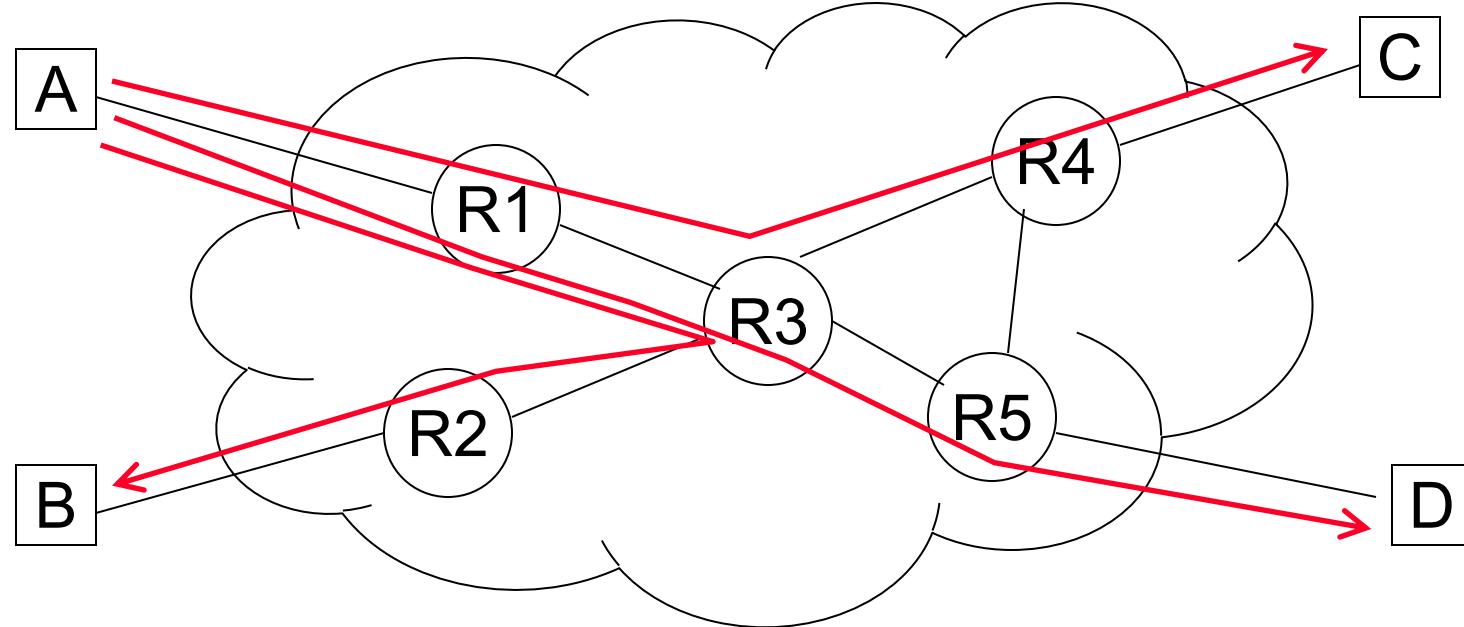
- The original paper [Birrell and Nelson'84]:
 - “Implementing Remote Procedure Calls”, in ACM Transactions on Computer Systems, Feb 1984.
- SUN RPC:
 - <http://www.rfc-editor.org/rfc/rfc1057.txt>
- XML-RPC:
 - <http://www.xmlrpc.com/>

Motivation: Multicast

- Single sender sends the same data to n receivers
- Use n point-to-point messages
 - Overload the sender – throughput limited by
server's bandwidth / n
- Multicast: Have the clients help forward the message to others

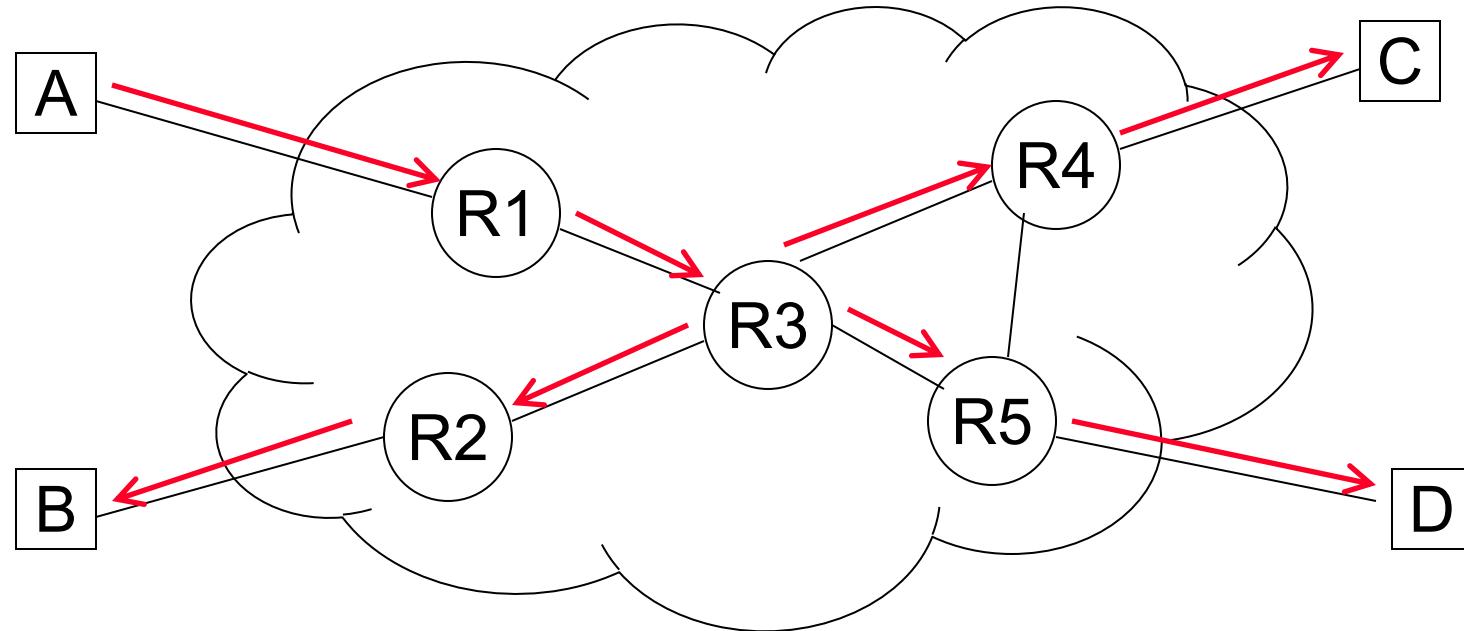
Without Multicast

- A uses its outgoing link 3 times
- Edge R1→R3 is used 3 times as well



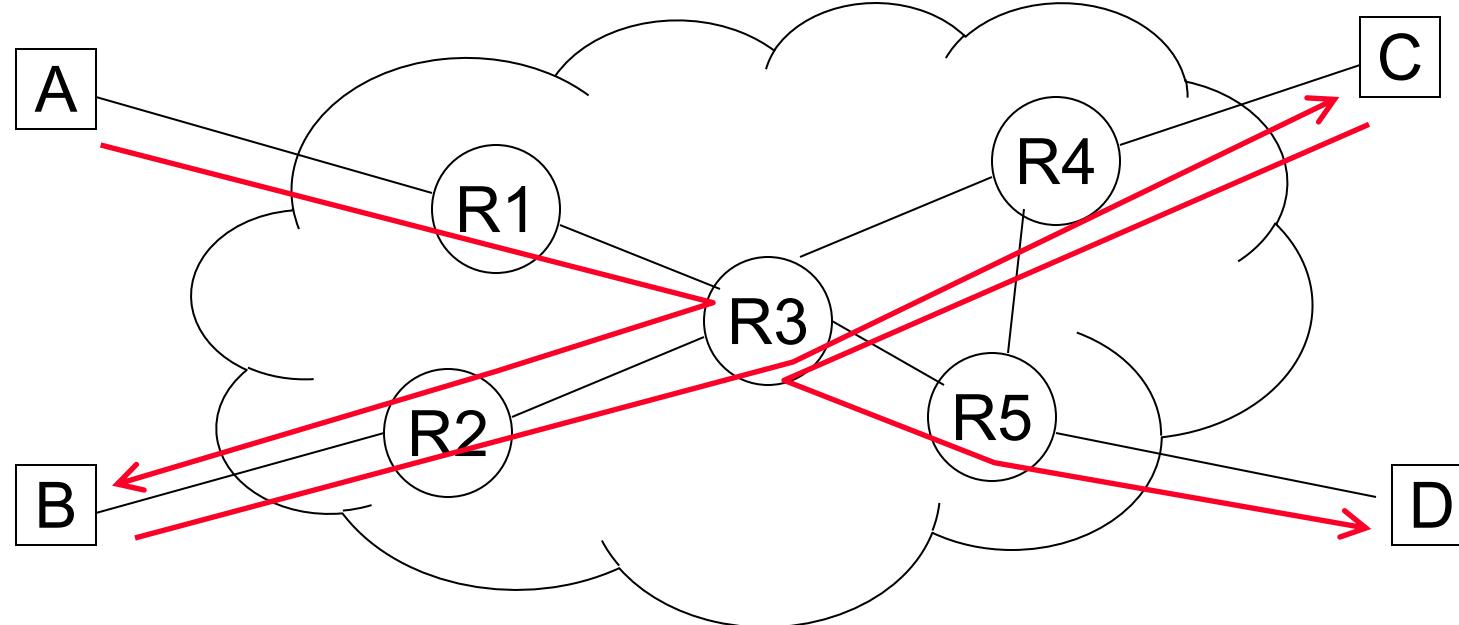
IP-Multicast

- A networking topic
- Very efficient – message never traverse same link twice
- Need support for routers – ISPs hate changes
- IP multicast was never hugely successful – a good technique does not always turn into products



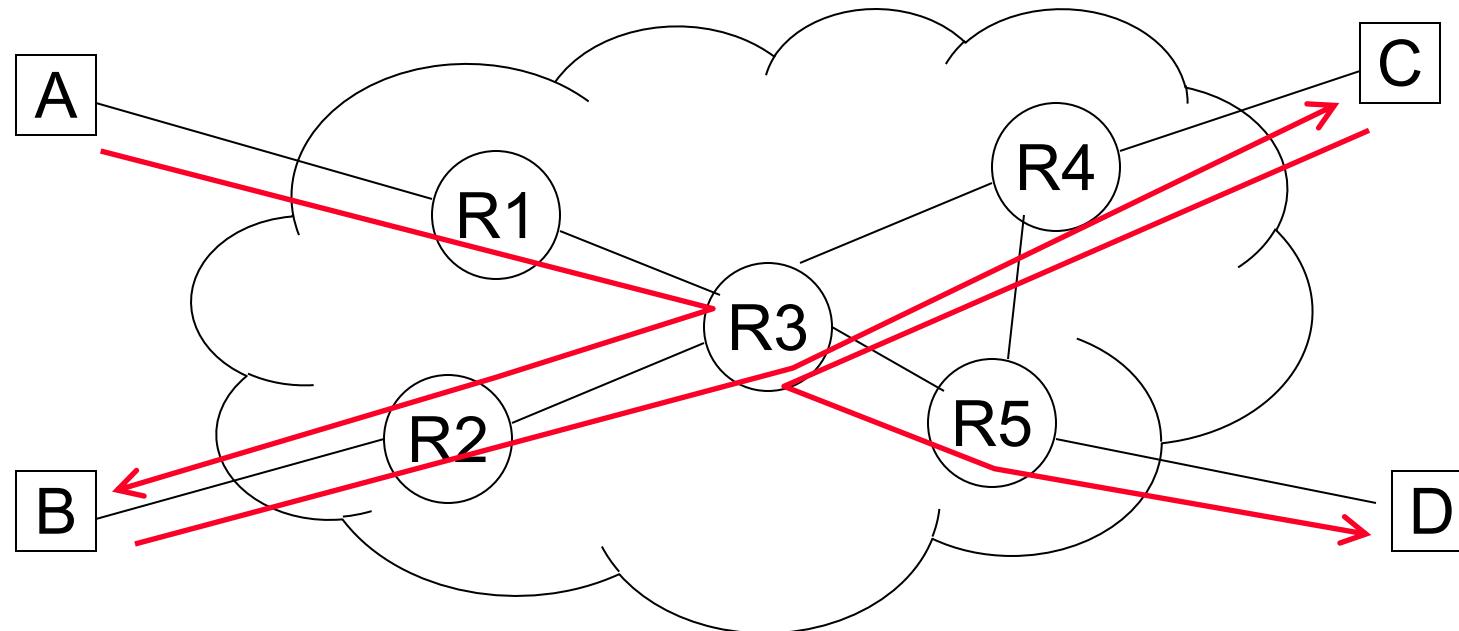
Application-Level Multicast

- Example: Server sends data to a small number of client, each client forwards the data to some other clients
- A distributed system topic



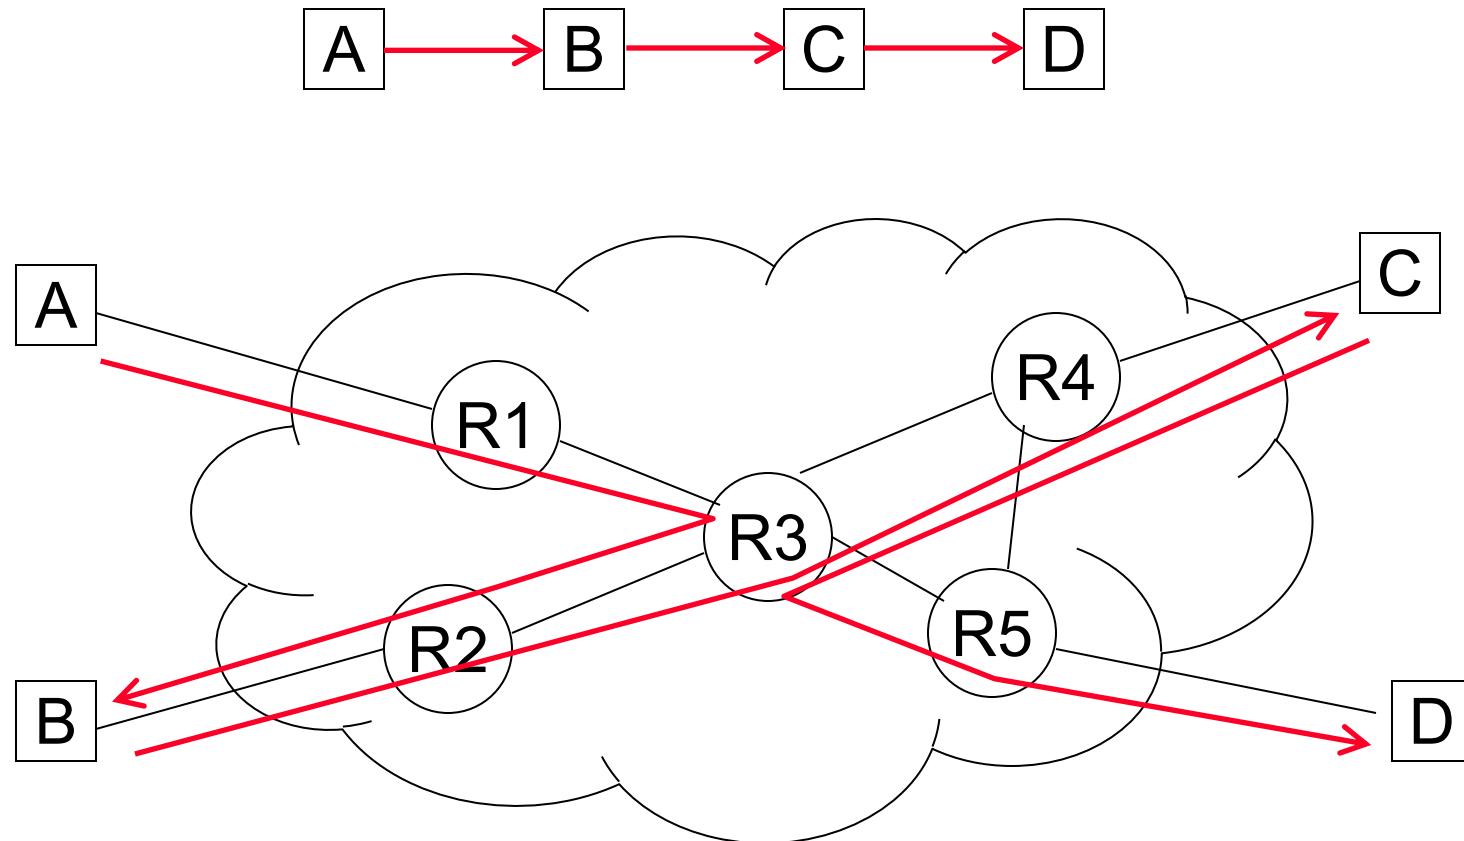
Application-Level Multicast

- Advantage:
 - Partly solves the problem with server bandwidth
 - No need to change routers
- Disadvantage:
 - A message may still traverse the same link twice
- Overall: Has been much more successful than IP multicast



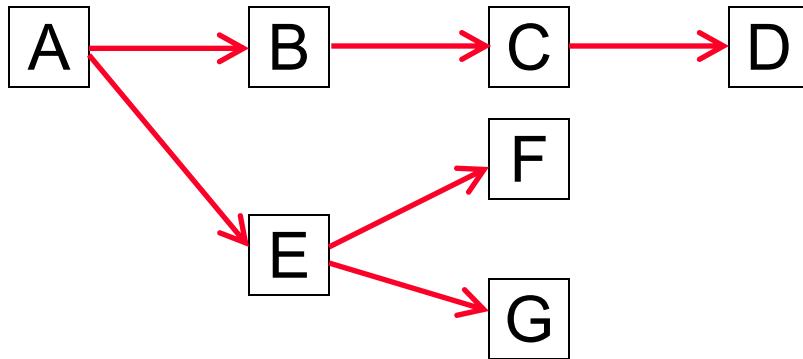
Designing Application-Level Multicast

- Multicast topology

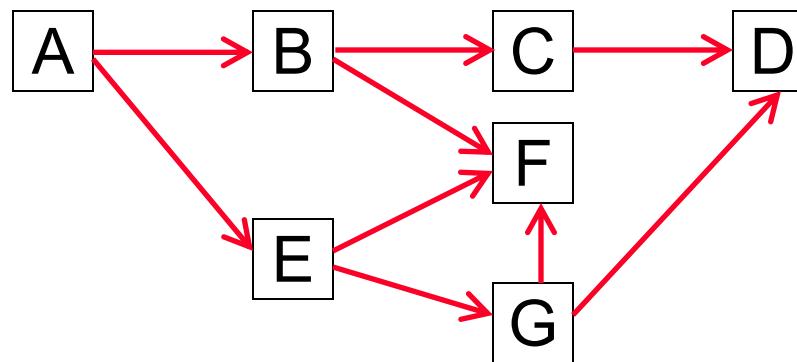


Multicast Topology

- Tree: Simple but not robust

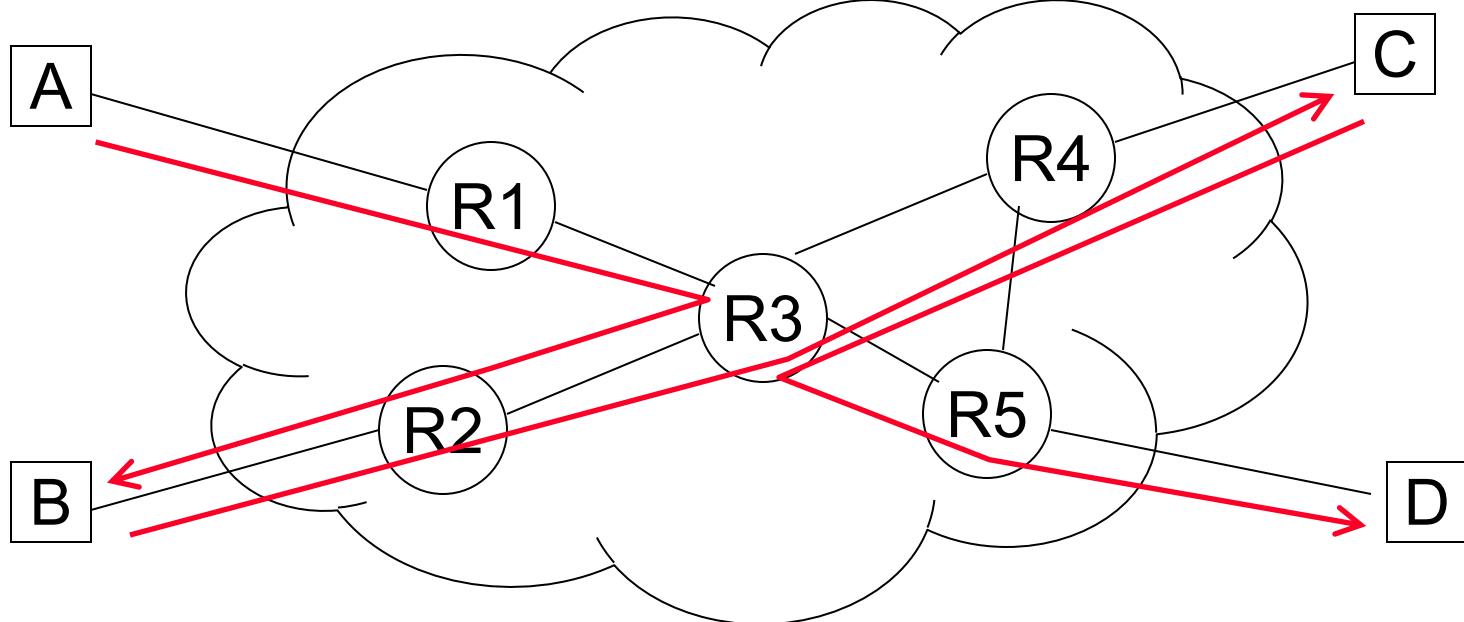


- Mesh: Robust but more complex – what data should I get from whom?



Designing Application-Level Multicast

- Quality of a multicast tree
 - **Link stress (2)**: maximum # messages traversing the same edge
 - **Stretch (12/4=3)**: Relative delay penalty given a pair of nodes
 - **Tree cost (12)**: Total weight of edges – minimum spanning tree optimizes for tree cost



Motivation: Gossip

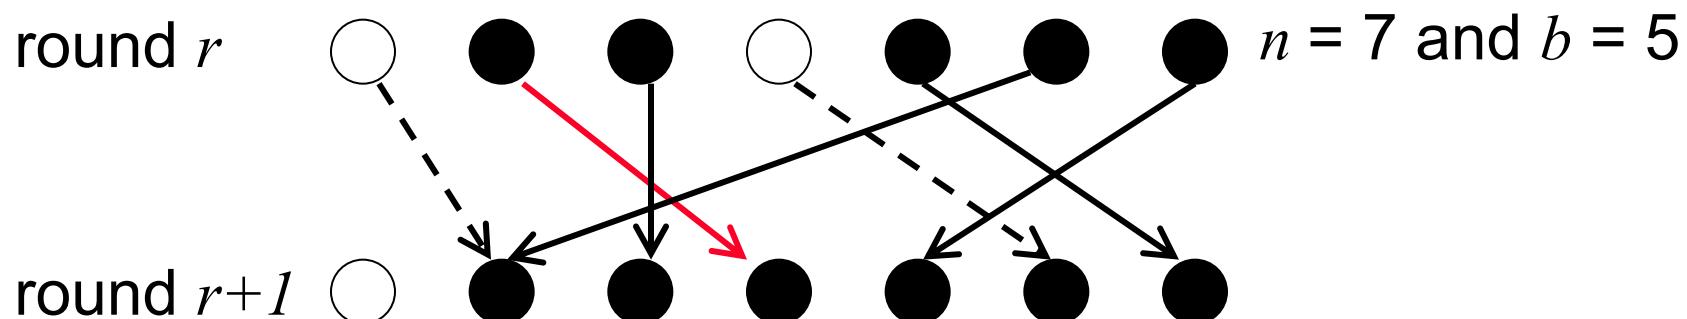
- Sender wants to send the same data to multiple receivers
 - Each node periodically send the data (if it has the data) to some other random node – usually called a **round**
- Advantage of gossip:
 - Reduces load of the sender
 - Extremely robust
 - Easy to verify from our everyday experience
 - Simple – no need to maintain a multicast topology
- Disadvantage – not as efficient as multicast

Designing a Gossip Protocol

- Each gossip operation is between two nodes
- How to pick the other node to gossip with?
 - Usually picked uniformly at random
 - How to pick random nodes without global knowledge?
 - Performance will be influenced by the choice
- How to gossip? If P picks Q
 - P may gossip news to Q (push)
 - P may request news from Q (pull)
 - P and Q may update each other (push+pull)

Performance of Push-based Gossiping

- P pushes message to Q
 - Effective at the beginning, but not toward the end
 - Chance of a black P picking a white Q decreases with time
 - n : # total nodes; b : # black nodes;
 - $\text{Prob}[\text{a black P picks a white Q}] = (n-b)/n = 2/7$
 - **But the number of black P's increases as well !**
(Description on textbook is not complete)



Push: Be More Rigorous

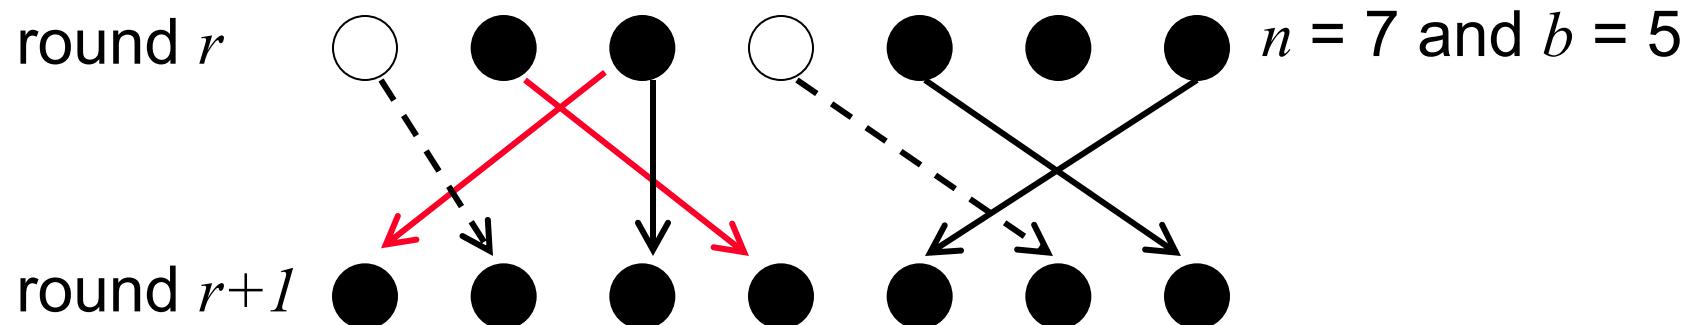
- Consider a white node Q at round r
 - Prob[Q remains white at round $r+1$] is $\left(1 - \frac{1}{n}\right)^b$
- Changing the last remaining white node to black (i.e., from $b = n-1$ to $b = n$)

$$1 - \left(1 - \frac{1}{n}\right)^{n-1} \approx 1 - \left(1 - \frac{1}{n}\right)^n \approx 1 - \frac{1}{e} \approx 0.63$$

- On expectation $1/0.63 = 1.6$ rounds needed
- Changing the first remaining white node to black (i.e., from $b = 1$ to $b = 2$)
 - On expectation $1/\left(1 - \frac{1}{n}\right) \approx 1$ round needed

Performance of Pull-based Gossiping

- P pushes message to Q
 - Not effective at the beginning, but effective toward the end
 - Chance of a white P picking a black Q increase with time
 - n : # total nodes; b : # black nodes;
 - $\text{Prob}[\text{a white P picks a black Q}] = b/n = 5/7$



Pull: Be More Rigorous

- Changing the first remaining white node to black (i.e., from $b = 1$ to $b = 2$)
 - On expectation 1.6 rounds needed – why?
- Changing the last remaining white node to black (i.e., from $b = n-1$ to $b = n$)
 - On expectation $1/\left(1 - \frac{1}{n}\right) \approx 1$ round needed
- One would naturally imagine that we want to use push+pull

Asymptotic Properties

- The asymptotic number of rounds needed to change all nodes to black is the same: $\theta(\log n)$
 - For push-only
 - For pull-only
 - For push+pull
- To prove this, consider two stages:
 - First stage: From $b = 1$ to $b = 0.1n$
 - Second stage: From $b = 0.1n$ to $b = n$

Second Stage

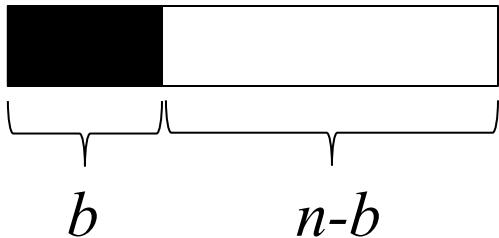
- At each round, imagine that the black nodes each throws a ball into a uniformly random bin (this is for push – pull is slightly different)
 - Total $0.1n$ balls thrown into n bins
 - The number of balls thrown at each round never below $0.1n$
- Coupon collection problem tells us:
 - On expectation needs $\theta(n \log n)$ balls
- Thus on expectation, # rounds needed is at most

$$\frac{\theta(n \log n)}{0.1n} = \theta(\log n)$$

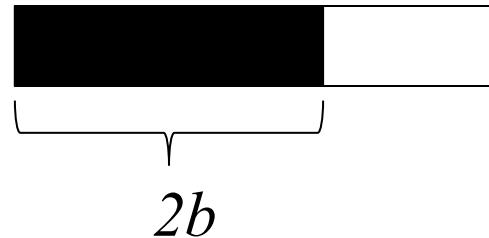
First Stage

- Define X_1 to be # rounds needed from $b = 1$ to $b = 2$
- Define X_2 to be # rounds needed from $b = 2$ to $b = 4$
- ...
- Can prove that $E[X_i] < 1.3$ for all i
 - Thus the first stage takes $\theta(\log n)$ rounds on expectation

Need b successes. Each ball has a success probability of at least $(n-b)/n \geq 0.8$.



Need on expectation $b/0.8 < 1.3b$ balls. Each round has at least b balls.
Need on expectation < 1.3 rounds.



History Readings (Non-compulsory)

- The original paper on gossiping:
 - “Epidemic Algorithms for Replicated Database Maintenance” in PODC’87
 - <http://portal.acm.org/citation.cfm?id=41841>

More Applications of Gossiping

- Gossiping not only good for propagating messages
- Another interesting application: Information aggregation

- Count the total number of nodes in the system
 - A distinguished node 0 starts from a value $x_0 \square 1$
 - Others have $x_i \square 0$

- Each node i periodically picks another node j
 - Node i : new $x_i \square \lfloor x_i \square x_j \rfloor / 2$
 - Node j : new $x_j \square \lfloor x_i \square x_j \rfloor / 2$

Counting the Number of Nodes

- Eventually all nodes will have the same x_i values
- Property from how we update x_i : Mass preserving
 - The sum of all x_i 's are always 1
- Property from gossiping:
 - All x_i 's will eventually be the same
 - Takes logarithmic number of rounds to achieve this (a much harder proof!)
- Number of nodes = $\frac{1}{x_i}$

Computing Average

- Each node has a value (e.g., available disk space)
 - We want to know the total available disk space
- Same as before but set x_i to be the amount of disk space available
- What if we want to compute sum?
- Drawback of this approach: Node leaves and node failures
 - Disrupts the mass preserving property – any way to patch?

Another Way of Computing Count

- Each node picks a random number between [0,1]
 - Use gossiping to find out the minimum
 - (The textbook uses maximum, which is the same – but conventionally people use minimum)
 - The larger the number, the smaller the minimum
- How to do sum and average?
- Advantage: More robust against failures
- Other in-depth differences with the previous approach exist (e.g., specific forms of the asymptotic properties)

History Readings (Non-compulsory)

- The beautiful push-sum algorithm
 - “Gossip-based Aggregation in Large Dynamic Network”
 - <http://portal.acm.org/citation.cfm?id=1082470>

Summary

- Review (Overview) of the OSI model
- Remote Procedure Call / Remote Method Invocation
 - If you don't know Java, real some toturials on Java
- Multicast
- Gossiping

Mandatory Homework For This Week

- Try out RMI by following tutorial at
 - Canvas\Files\AssignmentOne\RMI-tutorial.txt
 - You should be able to smoothly run and completely understand this example, otherwise you will have trouble with Assignment One
 - Email me or the TA (Yu Xiaoxue, e0323328@u.nus.edu) if you need help on this homework --- everyone should complete this homework
- Resources:
 - Canvas\Discussions -- You are encouraged to discussed with other students about this homework
- Security warning – kill your processes after running:
 - Do not leave rmiregistry running forever
 - Do not leave your server code running forever