

CS5223

Distributed Systems

Lecture 4: Naming

Instructor: YU Haifeng

Today's Roadmap

- Chapter 5 of textbook
 - P2P and DHT is a important topic and will be covered in a separate lecture
- Basic concepts in naming
- Simple naming mechanisms
- Domain name service

Motivation for Naming

- Example: We want to visit `www.nus.edu.sg`
 - The web server is located in some room that we don't know
 - There may be multiple web servers
 - ISP's may want to serve the content from their own machines – Akamai
 - `www.nus.edu.sg` → ip address → some physical machine (with specific MAC address)
- Example: Distributed file system
 - We know the file name (and full path), need to find out where the file is (i.e., on which machine, which disk sector, etc)
 - The file may migrate from one machine to another without changing its name

Basic Concepts in Naming

- General definition for naming:
 - We have a “name” and we want to find/locate the “entity”
 - Definitions in the textbook is unnecessarily convoluted, and sometimes circular – please ignore those
 - A lot of things can be called “naming” – just a definition issue – no need to draw a black-and-white boundary
- Name space: The set of all possible names in the given context
- Naming system: The system that enables you to find the “entity”
 - Name resolution
 - The naming system can be either non-distributed or distributed
 - Same tradeoff as before...

Different Kinds of Names

- Every name is a bit string
 - Some are human-friendly
- Hierarchical versus flat
- Alias
 - Multiple names for the same entity
 - Similar as symbolic links

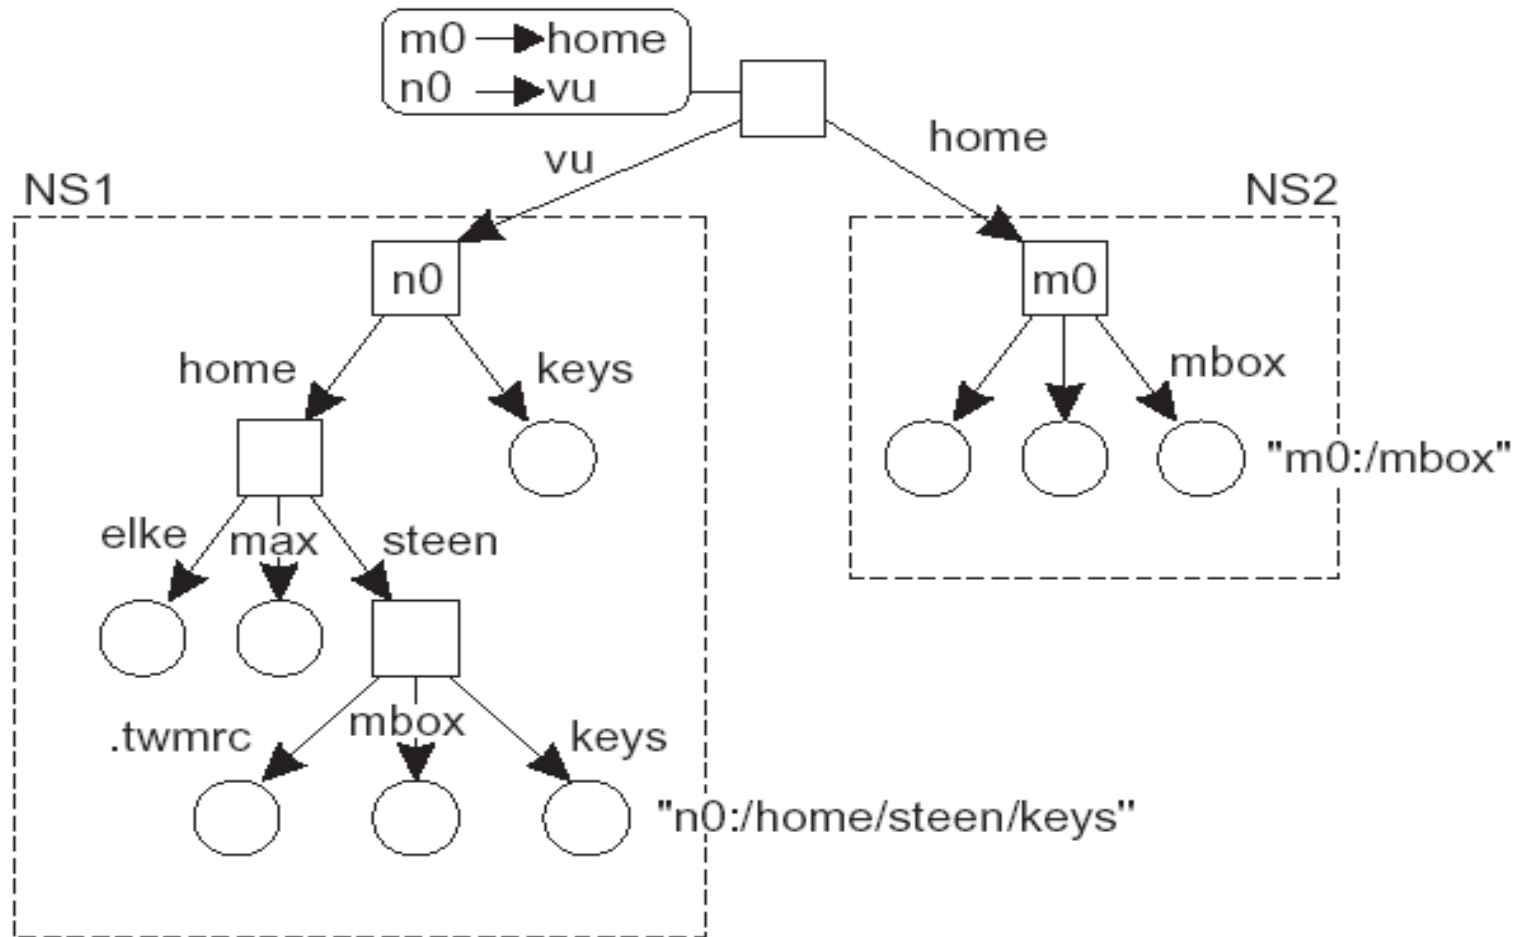
Merging Name Spaces

- We have two different name spaces
 - Now we want to define a unified name space
 - Name space merging
 - Generalizes to multiple name spaces...how?
- Example: Merging two network file systems
- Example: Two users want to give each other access to files in their home directories
- Main issue to take care of: There can be name collisions

Solution 1: New Root Node

- Merge by adding a new root node and make existing root nodes its children
- In principle, you always have to start in the new root
 - But you can avoid this by including the identifier of node from where resolution should start
 - Example, ~/XXX
- Problem: existing names need to be changed

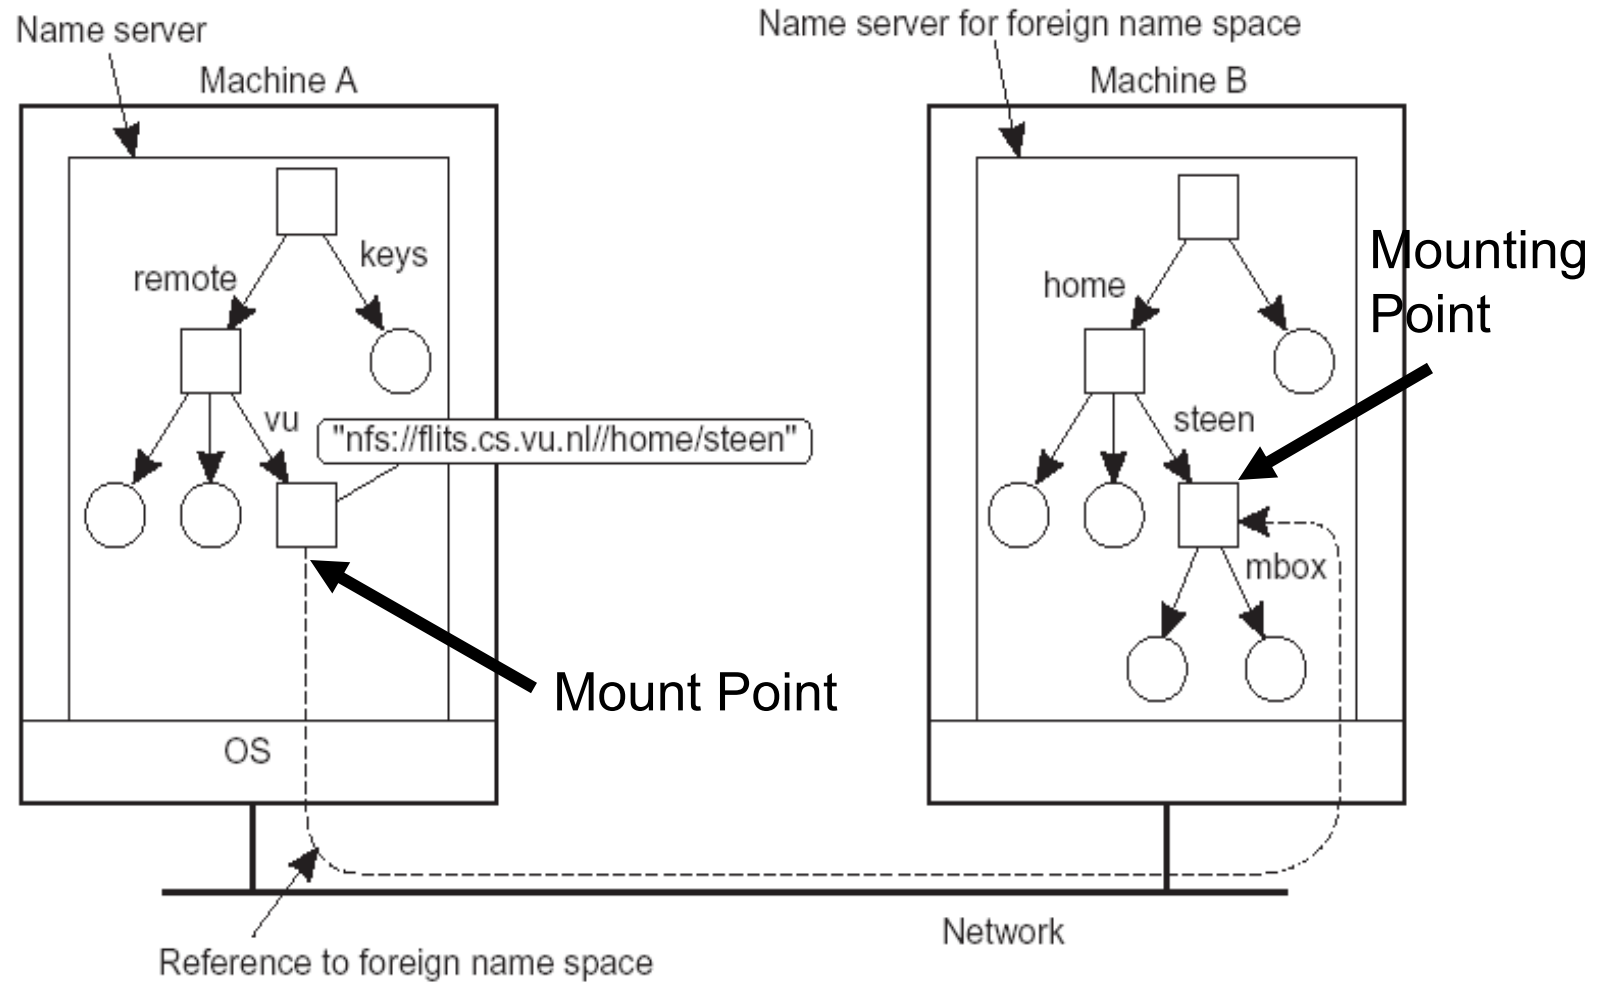
Example: New Root Node



Solution 2: Mounting

- Introduce nodes that contain the name of a node in a “foreign” name space
- Mount point: (Directory) node in naming graph that refers to other naming graph
- Mounting point: (Directory) node in other naming graph that is referred to

Example of Mounting



Simple Naming: Broadcast

- The “naming system” may or may not be a stand-alone software module
- Broadcast to resolve a name
 - Used in early day of computing
 - Simple is good!
- **Drawback: performance**
 - Alleviated but not fully addressed by using multicast

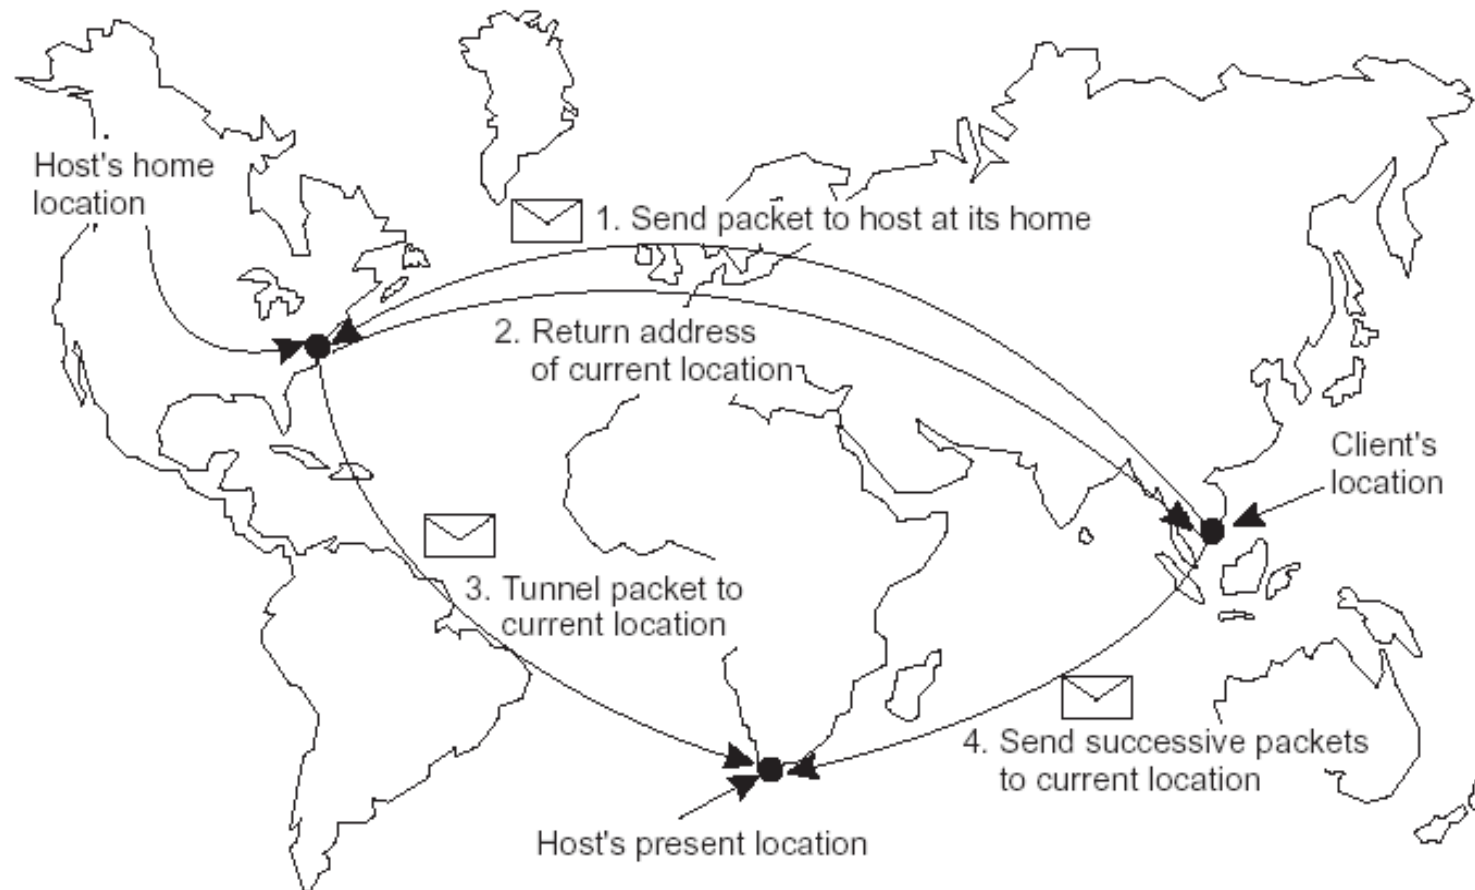
Simple Naming: Forwarding Pointers

- Forwarding pointers: Same idea as mail forwarding
- The machine A with ip_addr1 was initially the web server for www.nus.edu.sg
 - We want to use another B machine now with ip_addr2 as the server
- A will forward the HTTP requests to B
- Or A can inform the client of B's address
- Drawback:
 - A needs to stay there forever
 - The forwarding chain can become long

Simple Naming: Home-Based Approach

- Each entity has a “home”
 - The entity will register its IP address (which may change) with its home
- To find the entity, we contact the home first
- Only minor difference from forwarding pointers
 - Maximum one-hop forwarding
- Drawback:
 - The home needs to be there forever (i.e., have a fixed IP address)

Example: Home-Based Approach



Domain Name System: DNS

- For looking-up IP addresses
 - Example: what ip address `www.nus.edu.sg` corresponds to
- DNS dates back to over 30 years ago
 - Replaces the original Internet naming scheme that keep everything in a master file (lack of scalability and freedom for each domain to assign internal names)
- Overall, DNS has been a very successful distributed system
 - Many of the early design decisions turn out to be critical later
 - Additional reading (non-compulsory): “Signposts in cyberspace: The Domain Name System and Internet Navigation” by R. Levien, National Academic Research Council, 2005
- We will study it as a case study...

DNS Name Space

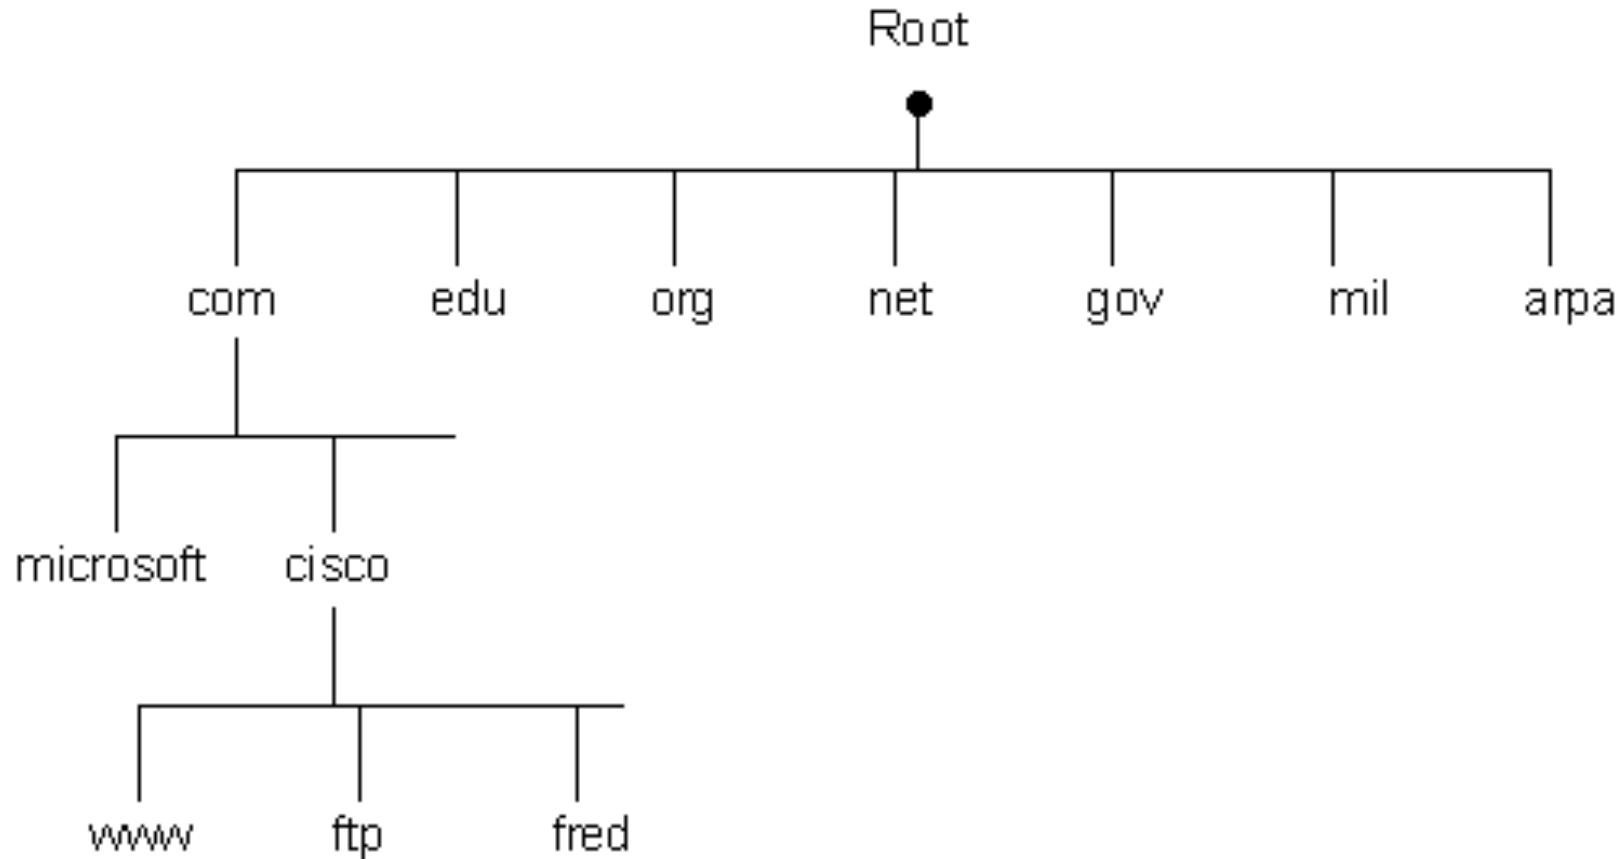
DNS name (max length 255)


www.comp.nus.edu.sg/~yuhf


label (max length 63)

- Top-level domains
 - com, edu, net, org, etc.
 - sg, cn, fr, etc.
- The machine corresponding to XXX.XXX.sg may or may not be in Singapore

The DNS Tree



(from <http://www.rhyshaden.com/dns.htm>)

Relative Names in DNS

- Not supported in DNS specification
- However:
 - Your local software may automatically append the omitted part
 - Example, you can type `ssh suna` on any SoC unix machine, and you will connect to `suna.comp.nus.edu.sg`
 - If `suna.comp.nus.edu.sg` does not exist, `ssh` will try to connect to “suna” (which does not exist either)
 - Sometimes the local software may try appending different suffixes

Domains in DNS

- A **domain** is a name space for which there is a single overall administrative authority to assignment the **DNS names**
 - Example, the collection of all XXX.nus.edu.sg is a domain
- A **domain's name** is the common suffix of all the DNS names in the collection
 - Example, nus.edu.sg is the domain's name
 - By definition, a domain's name is also a DNS name
- Potential confusion
 - A DNS name is sometime called **domain name**, which is different from domain's name defined above
 - We will only use DNS name and domain's name to avoid confusion

Resolving DNS Names: The Perspectives of a Client

- Web browser
 - Make query to DNS server (how to find the DNS server?) to resolve `www.nus.edu.sg` into an IP address
 - Send HTTP request to that IP address at port 80
- ftp, telnet, ssh, smtp(email) are all similar
- Reverse resolution:
 - Translate an IP address into a DNS name
 - Only for IP addresses within the local domain

Resolving a DNS Request

- The client needs to know the “local DNS server”
 - The IP address of this server is given to you by your ISP
 - Bootstrapping
- Client sends DNS query to local DNS server S
 - Containing the DNS name e.g., www.nus.edu.sg
- S checks whether it has either authoritative RR or cached RR for
 - www.nus.edu.sg
 - nus.edu.sg
 - sg

Root DNS servers

- If no, S will send quest to the root DNS server
 - How does S know where to find the root DNS server?
- Root servers with well-known IP addresses
- <https://www.iana.org/domains/root/servers>

Resolving a DNS Request

- The root DNS server sends back to S the IP address of the DNS server responsible for the domain “sg”
- S contact that DNS server
 - And then recursively the DNS server for “edu.sg”, “nus.edu.sg”, ...
- After a DNS name is resolved, the RR will be cached

Critical Design Decisions in DNS

- Delegation of responsibility
 - Allow individual domains to assign their own DNS names
 - Hierarchical naming structure
- Caching and replication
 - Controls the overhead of DNS lookups
 - Controls the load on the root DNS servers

Vulnerability of DNS

- The root servers – “single points” of failure
 - Target of DoS attacks
 - Already happened
- Alternatives:
 - Completely decentralized distributed naming system for the Internet
 - Already possible technically, but backward-compatibility issues need to be addressed

History Readings (Non-compulsory)

- DNS root servers under attack
 - http://en.wikipedia.org/wiki/DNS_Backbone_DDoS_Attacks

The Fundamental Need for Bootstrapping in Naming

- To resolve a name, you need to know at least one bootstrapping point
 - Broadcast: Bootstrapping point not needed (or every machine is a bootstrapping point?)
 - Forwarding pointers: The first machine you contact
 - Home-base approach: The home
 - DNS: Your local DNS server and the root DNS servers
- Called “closure mechanism” on the textbook
- Bootstrapping mechanism significantly influences the “goodness” of a naming system

Today's Summary

- Basic concepts in naming
- Simple naming mechanisms
- Domain name service